

Machine Learning

Daniel Mao

Contents

1	Formal Setup	1
1.1	Definitions	1
1.2	Empirical Risk Minimization (ERM)	1
2	Perceptron	3
2.1	The Perceptron Algorithm	3
3	Logistic Regression	5
3.1	Logit Transform	5
3.2	Hypothesis Representation	5
3.3	Decision Boundary	6
3.4	Maximum Likelihood	6
3.5	Cost Function	7
3.6	Optimization	8
3.7	Multiclass Classification	8
3.8	Introducing Non-linearity	8
3.9	Beyond Logistic Regression	9
4	Dimensionality Reduction	11
4.1	Motivations	11
4.2	Principal Component Analysis	11

Chapter 1

Formal Setup

1.1 Definitions

Definition (Error). Let \mathcal{X} denote the domain. Let \mathcal{D} be a probability distribution on X . Let f be the true labeling function on X . Let h be a classifier. We define the **error** of h , with respect to \mathcal{D} and f , denoted by $L_{\mathcal{D},f}$, to be a probability given by

$$L_{\mathcal{D},f}(h) := \mathcal{D}(\{x \in \mathcal{X} : h(x) \neq f(x)\}).$$

We are never sure about the true error. We only approximate the true error.

1.2 Empirical Risk Minimization (ERM)

Definition (Empirical Error). Let \mathcal{S} be a training data. Let h be a classifier. We define the **empirical error**, with respect to \mathcal{S} , denoted by $L_{\mathcal{S}}$, to be a proportion given by

$$L_{\mathcal{S}}(h) := \frac{|\{i : h(x_i) \neq y_i\}|}{|\mathcal{S}|}.$$

Chapter 2

Perceptron

2.1 The Perceptron Algorithm

Algorithm 1: The Perceptron Algorithm (Rosenblatt 1958)

Input: Dataset $\mathcal{D} = \{(\mathbf{x}_i, y_i) \in \mathbb{R}^d \times \{\pm 1\} : i = 1..n\}$, initialization $\mathbf{w} \in \mathbb{R}^d$ and $b \in \mathbb{R}$, threshold $\delta \geq 0$.

Output: Approximate solutions \mathbf{w} and b so that $\forall i \in \{1..n\}, y_i(\mathbf{w}^\top \mathbf{x}_i + b) > 0$.

```
1 while true do
2   receive training example index  $i \in \{1..n\}$ ;
3   if  $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \leq \delta$  then
4      $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$ ;
5      $b \leftarrow b + y_i$ ;
```

Remark. *We only perform updates when we make a mistake, and this is a necessary rule.*

Theorem 1. *Assume there exists some \mathbf{z} such that $A^\top \mathbf{z} > 0$, then the perceptron algorithm converges to some \mathbf{z}^* . If each column of A is selected indefinitely often, then $A^\top \mathbf{z}^* > \delta$.*

Corollary. *Let $\delta = 0$ and $\mathbf{z}_0 = \mathbf{0}$. Then the perceptron algorithm converges after at most $(R/\gamma)^2$ steps, where R and γ are dataset properties given by*

$$R := \|A\|_{2,\infty} = \max_i \|\mathbf{a}_i\|_2 \quad \text{and} \quad \gamma := \max_{\|\mathbf{z}\|_2 \leq 1} \min_i \langle \mathbf{z}, \mathbf{a}_i \rangle.$$

Note that the parameters R and γ are independent of the dataset size n .

Theorem 2. *The iterate $\mathbf{z} = (\mathbf{w}; b)$ of the perceptron algorithm is always bounded. In particular, if there is no separating hyperplane, then perceptron cycles.*

Chapter 3

Logistic Regression

3.1 Logit Transform

Definition (Logit Transform).

$$\log \frac{p(\mathbf{x}; \mathbf{w})}{1 - p(\mathbf{x}; \mathbf{w})} = \mathbf{w}^\top \mathbf{x}.$$

Or equivalently,

$$p(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})}.$$

This $p(\mathbf{x}; \mathbf{w})$ is the confidence in the prediction.

3.2 Hypothesis Representation

Logistic Regression Model

Want $0 \leq h_\theta(x) \leq 1$.

$$h_\theta(x) = g(\theta^\top x)$$

where

$$g(z) = \frac{1}{1 + e^{-z}}.$$

So

$$h_\theta(x) = \frac{1}{1 + e^{-\theta^\top x}}.$$

Sigmoid function (logistic function)

$$g(z) = \frac{1}{1 + e^{-z}}.$$

Properties:

- $g(0) = 0.5$.
- $\lim_{z \rightarrow +\infty} g(z) = 1$.
- $\lim_{z \rightarrow -\infty} g(z) = 0$.

Interpretation of the Hypothesis Output

$h_\theta(x)$ is the probability that $y = 1$ on input x . i.e. $h_\theta(x) = P(y = 1 \mid x; \theta)$, the probability that $y = 1$, given x , parameterized by θ .

Since y is either 0 or 1,

$$P(y = 0 \mid x; \theta) + P(y = 1 \mid x; \theta) = 1, \text{ or}$$

$$P(y = 0 \mid x; \theta) = 1 - P(y = 1 \mid x; \theta).$$

3.3 Decision Boundary

Proposition 3.3.1. *The decision boundary \mathcal{D} of logistic regression with parameter $\mathbf{w} \in \mathbb{R}^n$ is*

$$\mathcal{D} = \{x \in \mathbb{R}^n : \mathbf{w}^\top \mathbf{x} = 0\}.$$

Proof. Predict $y = 1$ if $h_{\mathbf{w}}(x) \geq 0.5$ and predict $y = 0$ if $h_{\mathbf{w}}(x) < 0.5$. Notice

$$h_{\mathbf{w}}(x) = 0.5 \iff \mathbf{w}^\top \mathbf{x} = 0.$$

So the model predicts $\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x})$ but with confidence $p(\mathbf{x}; \mathbf{w})$. ■

Example 3.3.1. Consider $\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$. Predict “ $y = 1$ ” if $\theta^\top \mathbf{x} = -3 + x_1 + x_2 \geq 0$. That is, $x_2 \geq 3 - x_1$. This gives a “half space” in \mathbb{R}^2 . The line $x_1 + x_2 = 3$ separates \mathbb{R}^2 into a region where the model predicts “ $y = 0$ ” and a region where the model predicts “ $y = 1$ ”.

3.4 Maximum Likelihood

Proposition 3.4.1. *The maximum likelihood function is*

$$\mathcal{L} = \prod_{i=1}^n p(\mathbf{x}_i; \mathbf{w})^{y_i} (1 - p(\mathbf{x}_i; \mathbf{w}))^{1-y_i}.$$

3.5 Cost Function

Training set: $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$. $x_0 = 1$. $x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$. $y \in \{0, 1\}$.

If

$$\text{cost}(h_\theta(x^{(i)}, y^{(i)})) = \frac{1}{2}(h_\theta(x^{(i)}) - y^{(i)})^2$$

as in linear regression, then the cost function is non-convex.

Logistic Regression Cost Function

$$\text{cost}(h_\theta(x), y) := \begin{cases} -\log(h_\theta(x)), & \text{if } y = 1 \\ -\log(1 - h_\theta(x)), & \text{if } y = 0. \end{cases}$$

Property:

- If $y = 1$ and $h_\theta(x) \uparrow 1$, then $\text{cost}(h_\theta(x), y) \approx 0$.
- If $y = 1$ and $h_\theta(x) \downarrow 0$, then $\text{cost}(h_\theta(x), y) \rightarrow +\infty$.
- Similar for $y = 0$.

Captures intuition that if $h_\theta(x) = 0$ (the model predicts $P(y = 1 \mid x; \theta) = 0$), but $y = 1$, we'll penalize the learning algorithm by a very large cost.

Cost Function for the Training Set

$$J(\theta) := \frac{1}{m} \sum_{i=1}^m \text{cost}(h_\theta(x^{(i)}, y^{(i)}))$$

where cost is the function given above.

Rewrite the cost function:

$$\text{cost}(h_\theta(x), y) = -y \log(h_\theta(x)) - (1 - y) \log(1 - h_\theta(x)).$$

So

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)})) \right].$$

To fit parameter θ : $\min_\theta J(\theta)$.

3.6 Optimization

Gradient Descent

$$\frac{\partial}{\partial \theta_j} J(\theta) = \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}.$$

This looks identical to linear regression.

Other Optimization Algorithms

- Conjugate Gradient
- BFGS
- L-BFGS

Advantages:

- No need to pick learning rate α .
- Often faster than gradient descent.

3.7 Multiclass Classification

One-vs-rest

$$h_{\theta}^{(i)}(x) = P(y = i \mid x; \theta) \quad i = 1, 2, 3.$$

Train a logistic regression classifier for each i to predict the probability that $y = i$. On a new input x , pick class $i = \operatorname{argmax}_i h_{\theta}^{(i)}(x)$.

Softmax

$$\mathbb{P}(Y = k \mid \mathbf{x}; \mathbf{W}) = \frac{\exp(\mathbf{w}_k^{\top} \mathbf{x})}{\sum_{l=1}^C \exp(\mathbf{w}_l^{\top} \mathbf{x})}.$$

3.8 Introducing Non-linearity

Consider

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

and $\theta = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$. Then the model predicts “ $y = 1$ ” if $-1 + x_1^2 + x_2^2 \geq 0$. That is, if $x_1^2 + x_2^2 \geq 1$.

3.9 Beyond Logistic Regression

The logit transform is just one choice among the others. What we need is a function Q from $[0, 1]$ onto \mathbb{R} . The inverse of any cumulative distribution function can be used as the Q function. e.g. the probit regression which uses the inverse of the CDF of a standard normal distribution.

Definition (Logistic Distribution). *We define the **logistic distribution** to be a probability distribution with CDF given by*

$$F(x; \mu; s) := \frac{1}{1 + \exp(-\frac{x-\mu}{s})}.$$

Proposition 3.9.1 (Mean of Logistic Distribution). *Let $X \sim F(x; \mu; s)$. Then*

$$\mathbb{E}[X] = \mu.$$

Proposition 3.9.2. *Let $X \sim F(x; \mu; s)$. Then*

$$\text{var}[X] = \frac{s^2 \pi^2}{3}.$$

Chapter 4

Dimensionality Reduction

4.1 Motivations

- Data compression.
- Data visualization.

4.2 Principal Component Analysis

PCA is trying to find a lower dimensional surface onto which to project the data so as to minimize squared projection error.

Principal component analysis vs. linear regression:

- Linear regression tries to minimize residuals. PCA tries to minimize projection error.
- In linear regression, we try to predict one feature with other features. In PCA, all features are equivalent.

Algorithm: (from n dimensions to k dimensions)

- Preprocessing: mean normalization and feature scaling (standardization).
- Compute covariance matrix

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)})(x^{(i)})^T.$$

- Compute the eigenvalues of Σ .

$$\Sigma = U S V.$$

Let U_{reduce} denote the first k columns of U . Then $z = U_{reduce}^T x$.

Choosing k :

- Average squared projection error:

$$\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{approx}^{(i)}\|^2.$$

- Total variation in the data:

$$\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2.$$

- Typically, we choose k to be the smallest value so that

$$\frac{\text{average squared projection error}}{\text{total variation}} < 0.01,$$

i.e. 99% of variance is retained.

- The above ratio can be computed in an easier way. Say

$$S = \begin{bmatrix} s_{11} & & O \\ & \ddots & \\ O & & s_{nn} \end{bmatrix}.$$

Then

$$\frac{\text{average squared projection error}}{\text{total variation}} = 1 - \frac{\sum_{i=1}^k s_{ii}}{\sum_{i=1}^n s_{ii}}.$$

So we can start with $k = 1$ and increase k and check whether

$$\frac{\sum_{i=1}^k s_{ii}}{\sum_{i=1}^n s_{ii}} > 0.99.$$